

Ali Al Kadhim and Harrison B. Prosper Department of Physics, Florida State University

```
import numpy as np; import pandas as pd
import scipy as sp; import scipy.stats as st
import torch; import torch.nn as nn
from numba import njit
from sklearn.preprocessing import StandardScaler; from
sklearn.model_selection import train_test_split
import matplotlib as mp; import matplotlib.pyplot as plt;
#reset matplotlib stle/parameters
import matplotlib as mpl
mpl.rcParams.update(mpl.rcParamsDefault)
plt.style.use('seaborn-deep')
mp.rcParams['agg.path.chunksize'] = 10000
font_legend = 15; font_axes=15
# %matplotlib inline
import copy; import sys; import os
from IPython.display import Image, display
import optuna
#sometimes jupyter doesnt initialize MathJax automatically for latex,
so do this
import ipywidgets as wid; wid.HTMLMath('$\LaTeX$')

{"version_major":2,"version_minor":0,"model_id":"3029f222dd247b582617
2e9f69ad1be"}
```

Table of Contents

Background and Problem Statement

In particle physics, the most important experiment is a counting experiment, represented by a Poisson probability model. Suppose we have a count experiment, we observe in each bin (or channel) k a count n_k . Given N total channels, the probability of obtaining the observed result is given by the Poisson

$$\prod_{k=1}^N \frac{e^{-(\epsilon_k \sigma + b_k)} (\epsilon_k \sigma + b_k)^{n_k}}{n_k!}$$

Where σ , the cross section (the parameter of interest), b_k is expected background for the k th channel (the nuisance parameter). ϵ_k is the "acceptance parameter", for the k th channel, which is typically a product of the detector efficiency, branching fraction, and luminosity.

This is the prototype of many statistical models in astronomy and particle physics in which data are binned and the count in each bin consist a priori of the sum of counts from signal and background. Written in terms of the simplified 2-parameter model, the expected count in each bin k takes the form

$$\begin{equation} n_{\text{exp}} = \theta + \nu \end{equation}$$

Where n_{exp} is the expected signal count, θ is the unknown mean count, which is the parameter of interest (the cross section), and ν is the background unknown mean count, which is the nuisance parameter.

The two-parameter problem

the conditional probability of observing n signal counts and m background counts in a single channel is given by

$$P(n, m \mid \theta, \nu) = L(\theta, \nu) = \frac{e^{-(\theta+\nu)} (\theta+\nu)^n}{n!} \frac{e^{-\nu} \nu^m}{m!}$$

Where, once the counts n and m have been observed becomes the likelihood $L(\theta, \nu)$. The objective now is to derive confidence intervals or limits for the parameter of interest θ . What makes this challenging is the presence of nuisance parameter ν . In the Bayesian approach, nuisance parameters are assigned priors $\pi(\nu) d\nu$ and are integrated out in order to arrive at the posterior for θ in a process called marginalization. Frequentist methods deal with nuisance parameters by profiling them the likelihood in order to arrive at their MLEs.

The typical way to arrive at confidence intervals that guarantee coverage is the Neyman-Pearson construction. However, it does not deal with the existence of nuisance parameters....

the Poisson mean $\theta + \nu$, and the probability model is given by

$$p(N, M \mid \theta, \nu) = \text{Poiss}(N; \theta + \nu) \text{Poiss}(M; \nu) = L(\theta, \nu) = \frac{e^{-(\theta+\nu)} (\theta+\nu)^N}{N!} \frac{e^{-\nu} \nu^M}{M!}$$

Where ν is the nuisance parameter, θ is the parameter of interest, for which we want to estimate confidence intervals or (upper) limits.

Observed Data:

- N (observed counts for signal)
- M (observed counts for background)

Parameters:

- θ : parameter of interest, proportional σ in HEP (unknown signal mean)
- ν : nuisance parameter (unknown background mean)

Auxiliary (simulated) Data (simulated on-the-fly for each observation)

- n : expected signal count
- m : expected background count

The standard procedure for removal of nuisance parameters is Bayesian marginalization, but in this study we adopt LFI with frequentest methods.

In this notebook we generate data comprising the quadruplets $(Z_i, \theta_i, \nu_i, N_i, M_i)$ where

$$\begin{aligned} \theta &\sim \text{uniform}(0, 20), \nu \sim \text{uniform}(0, 20), n \sim \text{poisson}(\theta + \nu), m \sim \text{poisson}(\nu), N \sim \text{uniform}(0, 10), M \sim \text{uniform}(0, 10), \text{ and } Z = I[\lambda_{\text{gen}} \leq \lambda_D(D, \theta)] \end{aligned}$$

where the size of each of these samples is B' , I is the indicator function, and $\lambda_D(D, \theta) = \lambda(N, M, \theta)$ is our chosen test statistic

$$\lambda = -2 \log \frac{p(n, m \vee \theta, \hat{\nu}(\theta))}{p(n, m \vee \hat{\theta}, \hat{\nu}(\theta))} = -2 \log \frac{L_{\text{prof}}(n, m, \theta, \hat{\nu}(\theta))}{L_{\text{prof}}(n, m, \hat{\theta}_{\text{MLE}}, \hat{\nu}(\theta))},$$

where $L_{\text{prof}}(n, m, \theta, \hat{\nu}(\theta))$ is the profiled likelihood. The maximum likelihood estimate of $\hat{\theta}$ is given by

$$\hat{\theta}_{\text{MLE}} = n - m$$

.

Low-count data can sometimes yield spurious results, where the MLE of a parameter of interest θ , could yield a negative result. In the case that θ is the cross section, yielding a negative result is non-physical, which leads to the ad-hoc fix: taking ignoring the MLE solution and taking $\hat{\theta} = 0$ when $n < m$,

$$\hat{\theta}_{\text{non-MLE}} = \begin{cases} n - m & n > m \\ 0 & n \leq m \end{cases}$$

$$\log p(n, m \vee \theta, \nu) = -(\theta + \nu) + n \log(\theta + \nu) - \nu + m \log \nu + \text{constants}$$

The MLE $\hat{\nu}(\theta)$ is attained by solving $\frac{\partial \log p(n, m \vee \theta, \nu)}{\partial \nu} = 0$ leading to

$$\hat{\nu}(\theta) = (g + \sqrt{g^2 + 8m\theta})/4$$

where $g \equiv n + m - 2\theta$.

Where $L_{\text{prof}}(n, m, \theta, \hat{\nu}(\theta))$ is the profiled likelihood - that is - the likelihood function when the nuisance parameters are replaced by their maximum likelihood estimates (MLE) for a given value of the parameter of interest. The test statistic in Eq. (1) is used due to its well-known behavior of it converging to a χ_k^2 variate with k degrees of freedom, according to Wilk's Theorem, where k is the number of free parameters (parameters of interest after the nuisance parameters have been replaced by their MLEs).

Clearly, our choice of using the test statistic above is motivated by the [Neyman-Pearson Lemma](#) Instead of using different statistics for LFI as in the [Ann Lee paper](#), we use this one

since in Ann lee's paper, for example the ACORE statistic

$$\Lambda(D; \Theta_0) := \log \frac{\sup_{\theta_0 \in \Theta_0} \prod_{i=1}^n O(X_i^{\text{obs}}; \theta_0)}{\sup_{\theta \in \Theta} \prod_{i=1}^n O(X_i^{\text{obs}}; \theta)},$$

since the odds $\mathbb{O}(\mathbf{X}_i^{\text{obs}})$

are merely estimates of the likelihood, and not the likelihood itself, therefore Wilk's theorem is not guaranteed to work, even if we have a large sample size.

```
def mkdir(dir_):
    """make a directory without overwriting what's in it if it
    exists"""
    # assert isinstance(dir_, str)
    try:
        os.system('mkdir -p %s' % str(dir_))
    except Exception:
        pass
```

```
mkdir('utils')
```

```
# %writefile 'src/utils.py'
```

```
def import_base_stack():
    import numpy as np; import pandas as pd
    import scipy as sp; from numba import njit
```

```
@njit
def DR(s, theta):
    return sp.special.gammainc(s, theta)
```

```
@njit
def DL(s, theta):
    return 1 - sp.special.gammainc(s+1, theta)
```

```
def L_prof(n,m,theta):
    k=1
    k1 = k+1
    k2 = 0.5/k1
    g = n+m - k1*theta
    nu_hat = k2* (g+ np.sqrt(g*g +4*k1*m*theta))
    p1 = st.poisson.pmf(n, mu = theta + nu_hat)
    p2 = st.poisson.pmf(m, mu = k * nu_hat)

    return p1*p2
```

```
def theta_hat(n,m, MLE=True):
    theta_hat = n-m
```

```

if not MLE:
    theta_hat = theta_hat * (theta_hat > 0)
return theta_hat

```

```

# @njit
def lambda_test(theta, n, m, MLE=True):
    Ln = L_prof(n,m,theta)
    Ld = L_prof(n,m, theta_hat(n,m, MLE))
    lambda_ = -2*np.log(Ln/Ld)
    return np.array(lambda_)

```

Prior to reporting the results of our method, we validate that our method by comparing it to the well-known results of Wilk's Theorem. That is, the test statistic $\lambda(\theta, n, m, \nu(\theta))$ should be distributed as a χ_1^2 (a χ^2 distribution with a number of degrees of freedom equal to the number of free parameters left in our problem).

We test this theorem with our algorithm, but stopping at step 9. We then histogram the cumulative distribution function (CDF) of $\lambda(\theta)$ for a given (fixed) θ and ν , and compare it to the analytical CDF of a χ_1^2 distribution. The figure (below) shows that the results of our test statistic does indeed agree with what we expect from Wilk's Theorem.

```
chi2_exp_size=40
```

```

def run_sim(theta, nu, MLE, lambda_size):
    """Sample  $n \sim \text{Pois}(\theta + \nu)$ ,
         $m \sim \text{Pois}(\nu)$ ,
        and compute
         $\lambda(\theta, n, m)$ 

    return: (n, m, lambda_), where each are np arrays of length
    lambda_size
    """
    n = st.poisson.rvs(theta+nu, size=lambda_size)
    m = st.poisson.rvs(nu, size=lambda_size)
    lambda_ = lambda_test(theta, n, m, MLE=MLE)
    return (n, m, lambda_)

def run_sims(points, MLE):
    """
    Run an entire simulation, that is, generate n and m from
    run_sim above, and calculate lambda, for

    input: a tuple of (theta, nu) scalars

    Returns: df, lambda_results

    where lambda_results is a list of tuples
        (n, m, lambda_, theta, nu)
    and df is just a dataframe of [n,m,lambda,theta,nu]
    """

```

```

"""
lambda_results=[]
df=pd.DataFrame()
for p in points:
    theta, nu = p
    df['theta']=theta
    df['nu']=nu
    n, m, lambda_ = run_sim(theta, nu, MLE, lambda_size
=chi2_exp_size)
    df['n'] = n
    df['m'] = m
    df['lambda']=lambda_
    lambda_results.append((n, m, lambda_, theta, nu))

    print( '\n \n (theta, nu) = (%.f, %.f) \n ' % (theta, nu) )
    print(f'\t \t with associated n = {n}, \n \n \t \t m = {m}, \
n \n \t \t lambda = {lambda_}' )
    return df, lambda_results

def plot_one(lambda_, theta, nu, ax):
    """Histogram the CDF of lambda_t =
-2log(Lp(theta)/Lp(theta_hat)),
for a given (fixed) theta and nu.
Also, plot the actual CDF of a chi^2 distribution with 1 free
parameter
(since only theta is left after we profile nu) """
    fsize = 16; xmin= 0; xmax= 10
    ymin= 0; ymax= 1
    x_range = (xmin, xmax)
    y_range = (ymin, ymax)
    ax.set_xlim(x_range); ax.set_ylim(y_range)
    ax.set_xlabel(r'$\lambda$ (\theta, n, m, \hat{\nu}(\theta))
$', fontsize=fsize)
    ax.set_ylabel(r'cdf$(\lambda)$', fontsize=fsize)
    #####HISTOGRAM CDF OF LAMBDA#####
    ax.hist(lambda_, bins=5*xmax, range=x_range,
color=(0.8,0.8,0.9),
density=True, cumulative=True,
histtype='stepfilled', edgecolor='black', label=r'CDF$(\lambda)$')
    #####
    ##### HISTOGRAM CDF OF THE CHI2 OF OF X WITH 1 DOF
    #x is not theta, that's the whole point of Wilks thm, x is an
arbitrary RV
    x = np.arange(0, xmax, 0.2)
    y = st.chi2.cdf(x, 1)
    ax.plot(x, y, color='blue',
linewidth=2, label=r'CDF$(\chi^2_1)$')
    # annotate
    xwid = (xmax-xmin)/12

```

```

ywid = (ymax-ymin)/12
xpos = xmin + xwid/2
ypos = ymin + ywid*2
ax.text(xpos, ypos,
r'$ \theta = %d, \nu = %d$' % (theta, nu),
fontsize=ftsize)
ax.legend(loc='upper right',fontsize=15)

```

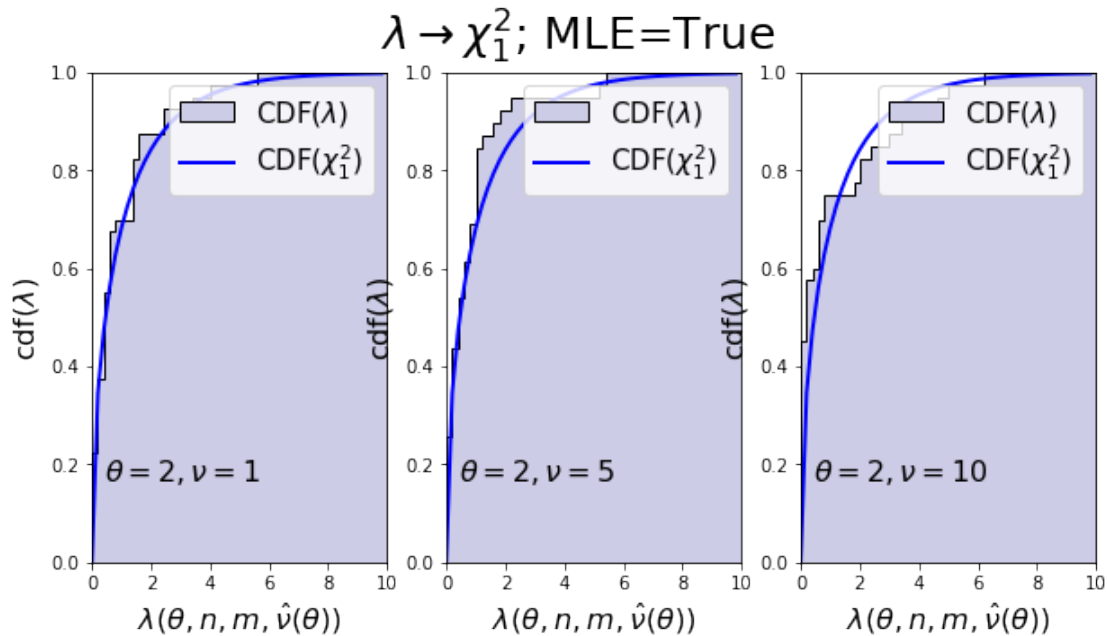
1. Generate one scalar θ and one scalar ν
2. Generate $\lambda(\theta, \nu)$ of size N . Observe that the distribution of this λ will approach the CDF of a $\chi^2_{dof}(x)$ of an RV of one dof (since θ is the only free parameter left), confirming Wilk's theorem
3. Observe that this test statistic will be dependent on the value of ν in the non-MLE case, which is not desirable since we want to be insensitive to nuisance parameters for maximal statistical power

```

#points=(theta,nu)
points_1 = (2, 1)
points_2 = (2, 5)
points_3 = (2, 10)
MLE=True
_, _, lambda_1 = run_sim(theta=points_1[0], nu=points_1[1], MLE=MLE,
lambda_size=chi2_exp_size)
_, _, lambda_2 = run_sim(theta=points_2[0], nu=points_2[1], MLE=MLE,
lambda_size=chi2_exp_size)
_, _, lambda_3 = run_sim(theta=points_3[0], nu=points_3[1], MLE=MLE,
lambda_size=chi2_exp_size)

fig, ax = plt.subplots(1,3, figsize=(10,5))
plot_one(lambda_1, points_1[0], points_1[1], ax[0])
plot_one(lambda_2, points_2[0], points_2[1], ax[1])
plot_one(lambda_3, points_3[0], points_3[1], ax[2])
fig.suptitle(r'$ \lambda \rightarrow \chi^2_1$; MLE=%s' % str(MLE),
fontsize=25); plt.show()


```



```
#points=(theta,nu)
points_1 = (2, 1)
points_2 = (2, 5)
points_3 = (2, 10)
MLE=False

_, _, lambda_1 = run_sim(theta=points_1[0], nu=points_1[1], MLE=MLE,
lambda_size=chi2_exp_size)
_, _, lambda_2 = run_sim(theta=points_2[0], nu=points_2[1], MLE=MLE,
lambda_size=chi2_exp_size)
_, _, lambda_3 = run_sim(theta=points_3[0], nu=points_3[1], MLE=MLE,
lambda_size=chi2_exp_size)

fig, ax = plt.subplots(1,3, figsize=(10,5))
plot_one(lambda_1, points_1[0], points_1[1], ax[0])
plot_one(lambda_2, points_2[0], points_2[1], ax[1])
plot_one(lambda_3, points_3[0], points_3[1], ax[2])
fig.suptitle(r'$\lambda \rightarrow \chi^2_1$; MLE=False',
fontsize=25); plt.show()
```

As shall be discussed later in this notebook, one way of calculating p -value for θ corresponding to an observed $\{N, M\}$ pair (which will still be dependent on ν) is by calculating the test statistic at the $\{N, M\}$ pair and calculating

$$p_{\theta}(\nu) = \text{Prob}(\lambda_{gen}(n, m; \theta, \nu) \leq \lambda_D(N, M; \theta, \nu)) = \int_{\lambda_D}^{\infty} f(\lambda_{gen}(n, m | \theta, \nu) | H_{null}) d\lambda_{gen},$$

where f is the PDF of λ . We see from above that the statistic λ is an approximate sampling distribution of the χ^2 PDF, proving [Wilk's theorem](#), and hence the p -value could also be calculated by

$$p_{\theta}(\nu) = \int_{\chi^2}^{\infty} f(z; k) dz$$

where k is the number of degrees of freedom (free parameters) which is 1 in the 2-parameter problem, and $f(z; k)$ is the χ^2 PDF. This p -value can then easily be computed (e.g. yielding $p = 1 - \alpha(1 - F_{\chi^2}[\theta, \nu])$, where F_{χ^2} is the cumulative χ^2 distribution function (as a function of θ and ν).

Generate 6 pairs (tuples) of (θ, ν) values

```
points = [(theta, nu) for theta, nu in
           zip(np.random.randint(low=1, high=4, size=3),
               np.random.randint(low=0, high=4, size=3))]
```

```
df, results = run_sims(points, MLE=True)
```

```

(theta, nu) = (2, 0)

with associated n = [2 5 4 3 1 3 2 0 1 1 1 2 1 1 2 4 1 1
2 3 1 1 3 3 1 4 3 2 1 3 4 2 2 3 0 3 2
0 3 2],

m = [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0],

lambda = [-0.          2.93147181  1.54517744  0.43279065
0.61370564  0.43279065
-0.          4.          0.61370564  0.61370564  0.61370564 -0.
0.61370564  0.61370564 -0.          1.54517744  0.61370564
0.61370564
-0.          0.43279065  0.61370564  0.61370564  0.43279065
0.43279065
0.61370564  1.54517744  0.43279065 -0.          0.61370564
0.43279065
1.54517744 -0.          -0.          0.43279065  4.
0.43279065
-0.          4.          0.43279065 -0.          ]

```

```

(theta, nu) = (2, 3)

with associated n = [5 9 5 9 3 6 4 4 6 2 8 4 4 5 2 7 5 3
5 4 6 5 5 8 7 4 6 2 5 2 4 7 4 5 7 4 4
3 9 8],

m = [3 4 6 5 4 1 5 3 5 4 6 8 5 5 3 2 4 5 3 5 5 3 3 0 4 9 2
4 4 3 5 2 0 1 4 2 3
3 2 2],

lambda = [-0.          0.7165121  0.80229178  0.29075642
1.2288131  1.47928694
0.97179616  0.1396362  0.09005222  2.51359335 -0.
2.96578989
0.97179616  0.39239727  1.66020193  1.07969347  0.10956358
1.92861676
-0.          0.97179616  0.09005222 -0.          -0.
7.09035489
0.09206738  3.7531945  0.52939123  2.51359335  0.10956358
1.66020193
0.97179616  1.07969347  1.54517744  0.74478362  0.09206738 -0.
0.1396362  0.63419698  2.52578322  1.75504047]

```

```

(theta, nu) = (1, 0)

```

```

with associated n = [2 0 0 1 1 1 1 2 0 2 2 0 0 3 2 3 1 3
1 3 0 1 1 3 0 1 1 4 1 0 0 2 3 2 1 0 2
1 0 0],

```

```

m = [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0],

```

```

lambda = [ 0.77258872  2.          2.          -0.
-0.          -0.
-0.          0.77258872  2.          0.77258872  0.77258872  2.
2.          2.15888308  0.77258872  2.15888308 -0.
2.15888308
-0.          2.15888308  2.          -0.          -0.
2.15888308
2.          -0.          -0.          3.54517744 -0.          2.
2.          0.77258872  2.15888308  0.77258872 -0.          2.
0.77258872 -0.          2.          2.          ]

```

```
df.head()
```

```

theta  nu  n  m  lambda
0      1   0  2  0  0.772589
1      1   0  0  0  2.000000
2      1   0  0  0  2.000000
3      1   0  1  0 -0.000000
4      1   0  1  0 -0.000000

```

```

n, m, lambda_, theta, nu = results[1]
print(f'results[1] = (n1, m1, lambda1, theta1, nu1) = ( \n {n}, \n
{m}, \n {lambda_}, \n {theta}, \n {nu})' )
print(f'''\n (n.shape, m.shape, lambda_.shape, theta.size, nu.size) =
({n.shape}, \t {m.shape}, \t {lambda_.shape}, \t {theta.size}, \
t {nu.size})''')

```

```

results[1] = (n1, m1, lambda1, theta1, nu1) = (
[5 9 5 9 3 6 4 4 6 2 8 4 4 5 2 7 5 3 5 4 6 5 5 8 7 4 6 2 5 2 4 7 4 5
7 4 4
3 9 8],
[3 4 6 5 4 1 5 3 5 4 6 8 5 5 3 2 4 5 3 5 5 3 3 0 4 9 2 4 4 3 5 2 0 1
4 2 3
3 2 2],
[-0.          0.7165121  0.80229178  0.29075642  1.2288131
1.47928694
0.97179616  0.1396362  0.09005222  2.51359335 -0.
2.96578989
0.97179616  0.39239727  1.66020193  1.07969347  0.10956358
1.92861676
-0.          0.97179616  0.09005222 -0.          -0.
7.09035489
0.09206738  3.7531945  0.52939123  2.51359335  0.10956358

```

```

1.66020193
  0.97179616  1.07969347  1.54517744  0.74478362  0.09206738 -0.
  0.1396362   0.63419698  2.52578322  1.75504047],
2,
3)

```

```

(n.shape, m.shape, lambda_.shape, theta.size, nu.size) =
  ((40,), (40,), (40,), 1, 1)

```

Generate Training data (or take a look at the saved training data)

We then generate training data where the number of training examples is B' according to Alg. 2 of Anne Lee et al. (shown below). The training data now has $\{\theta, \nu, N, M\}$ as training features and Z as the target. We then use Pytorch to build MLP regression model with average quadratic loss to estimate the distribution of $Z, E[Z \vee \theta, \nu]$, which according to Alg. 2 is the p-value.

```
algorithm2 = Image('images/Algorithm2.jpg'); display(algorithm2)
```

Algorithm 2 Estimate the p-values $p(D; \theta_0)$, given observed data D , for a level- α test of $H_{0, \theta_0} : \theta = \theta_0$ vs. $H_{1, \theta_0} : \theta \neq \theta_0$ for all $\theta_0 \in \Theta$ simultaneously.

Require: observed data D ; stochastic forward simulator F_θ ; sample size B' for p-value estimation; π_Θ (a fixed proposal distribution over the parameter space Θ); test statistic λ ; regression estimator m

Ensure: estimated p-value $\hat{p}(D; \theta)$ for all $\theta = \theta_0 \in \Theta$

```

1: Set  $\mathcal{T}' \leftarrow \emptyset$ 
2: for  $i$  in  $\{1, \dots, B'\}$  do
3:   Draw parameter  $\theta_i \sim \pi_\Theta$ 
4:   Draw sample  $\mathbf{X}_{i,1}, \dots, \mathbf{X}_{i,n} \stackrel{iid}{\sim} F_{\theta_i}$ 
5:   Compute test statistic  $\lambda_i \leftarrow \lambda((\mathbf{X}_{i,1}, \dots, \mathbf{X}_{i,n}); \theta_i)$ 
6:   Compute indicator  $Z_i \leftarrow \mathbb{I}(\lambda_i < \lambda(D; \theta_i))$ 
7:    $\mathcal{T}' \leftarrow \mathcal{T}' \cup \{(\theta_i, Z_i)\}$ 
8: end for
9: Use  $\mathcal{T}'$  to learn parametrized function  $\hat{p}(D; \theta) := \widehat{\mathbb{E}}[Z | \theta]$  via regression of  $Z$  on  $\theta$  using regression estimator  $m$ 
10: return  $\hat{p}(D; \theta_0)$ 

```

As we know, the p-value is the probability under the null hypothesis H_{null} (which is in this case parameterized by θ) of finding data of equal or greater *incompatibility* with the predictions of H_{null} . Therefore, in our case, the p-value under the null hypothesis (defined by θ) is given by

$$p_\theta(\nu) = \text{Prob}(\lambda_{gen}(n, m; \theta, \nu) \leq \lambda_D(N, M; \theta, \nu)) = \int_{\lambda_D}^{\infty} f(\lambda_{gen}(n, m | \theta, \nu) | H_{null}) d\lambda_{gen},$$

where f is the PDF of λ . In the strict frequentist approach, θ is rejected only if the p-value is less than the significance level of a hypothesis test α (i.e. accepted if $p_\theta(\nu) \leq \alpha$)

After calculating the p -value, one can then

In our study we approximate this integral as the histogram of θ weighted by Z divided by the histogram of θ , and denote this histogram as h . We compare this exact p -value with the outcome of our MLP, $\hat{p}(\theta, v, N, M)$ for the case where $\hat{\theta}$ is taken as the MLE in Fig. below we see almost perfect match.


```
train_df =
pd.read_csv('data/two_parameters_theta_0_20_1000k_Examples_MLE_True.csv')
train_df.describe()
```

	Unnamed: 0	Z	theta	nu
\count	1000000.000000	1000000.000000	1000000.000000	1000000.000000
mean	499999.500000	0.869337	10.002658	9.999062
std	288675.278933	0.337032	5.777129	5.772508
min	0.000000	0.000000	0.000013	0.000034
25%	249999.750000	1.000000	4.994221	4.996941
50%	499999.500000	1.000000	10.005912	9.997695
75%	749999.250000	1.000000	15.001766	14.996255
max	999999.000000	1.000000	19.999969	19.999976

	N	M
count	1000000.000000	1000000.000000
mean	5.001477	4.999767
std	2.581041	2.582273
min	1.000000	1.000000
25%	3.000000	3.000000
50%	5.000000	5.000000
75%	7.000000	7.000000
max	9.000000	9.000000

```
bins_ = 'auto'
fig, ax = plt.subplots(nrows=1, ncols=5, figsize=(20,10))
ax[0].hist(train_df['Z'], bins=bins_, label=r'$Z$')
ax[1].hist(train_df['theta'], bins=bins_, label=r'$\theta$')
ax[2].hist(train_df['nu'], bins=bins_, label=r'$\nu$')
ax[3].hist(train_df['N'], bins=bins_, label=r'$N$')
ax[4].hist(train_df['M'], bins=bins_, label=r'$M$')
```

```
[ax[i].legend(fontsize=font_legend) for i in range(5)]
plt.tight_layout(); plt.show()
```



Generate data according to:

$$\begin{aligned} \theta &\sim \text{uniform}(0, 20), \nu \sim \text{uniform}(0, 20), n \sim \text{poisson}(\theta + \nu), m \sim \text{poisson}(\nu), N \sim \text{uniform}(0, 10), M \sim \text{uniform}(0, 10), \\ Z &= I[\lambda_{\text{gen}} \leq \lambda_D(D, \theta)], \end{aligned}$$

```
thetaMin, thetaMax = 0, 20
numin, numax = 0, 20
Nmin, Nmax = 1, 10
Mmin, Mmax = 1, 10
```

```
def generate_training_data(Bprime, MLE, save_data=False):
    """Generate the training data, that is, features=[theta, nu, N,
M], targets=Z"""
    #sample theta and nu from uniform(0,20)
    theta = st.uniform.rvs(thetaMin, thetaMax, size=Bprime)
    # nu = st.uniform.rvs(nuMin, nuMax, size=Bprime)
    nu = st.uniform.rvs(numin, numax, size=Bprime)
    #n,m ~ F_{\theta, \nu}, ie our simulator. sample n from a Poisson
    with mean theta+nu
    n = st.poisson.rvs(theta+ nu, size=Bprime)
    #sample m from a poisson with mean nu
    m = st.poisson.rvs(nu, size=Bprime)
    #sample our observed counts (N,M), which take the place of D
    N = np.random.randint(Nmin, Nmax, size=Bprime)
    M = np.random.randint(Mmin, Mmax, size=Bprime)
    SUBSAMPLE=10
    print('n=', n[:SUBSAMPLE])
    print('m=', m[:SUBSAMPLE])
```

```

print('N=', N[:SUBSAMPLE])
print('M=', M[:SUBSAMPLE])
lambda_gen = lambda_test(theta, n, m, MLE)
print('lambda_gen= ', lambda_gen[:SUBSAMPLE])
lambda_D = lambda_test(theta, N, M, MLE)
print('lambda_D= ', lambda_D[:SUBSAMPLE])
#if lambda_gen <= lambda_D: Z=1, else Z=0
Z = (lambda_gen <= lambda_D).astype(np.int32)

data_2_param = {'Z' : Z, 'theta' : theta, 'nu': nu, 'N':N, 'M':M}

data_2_param = pd.DataFrame.from_dict(data_2_param)
if save_data:
    data_2_param.to_csv('data/two_parameters_theta_%s_%s_
%sk_Examples_MLE_%s.csv' %\
                        (str(thetaMin), str(thetaMax),
str(int(Bprime/1000)), str(MLE)) )

    print('\n')
    print(data_2_param.describe())
    return data_2_param

Train_data_MLE_True = generate_training_data(Bprime=1000000, MLE=True,
save_data=True)

n= [15 17  4 20  9 31 12 18 26 12]
m= [ 9  9  1  0  0  5  4  8 14  7]
N= [8 6 3 2 5 9 6 8 1 9]
M= [7 6 6 9 8 5 8 1 6 6]
lambda_gen= [0.22041922 0.59769146 0.70458812 0.15167097 1.59159138
3.72210763
1.44015718 0.21793792 0.27341912 2.61956417]
lambda_D= [ 3.19597616 1.32802053 1.95898303 35.34006927
5.32754169 6.69370749
12.47100404 0.04426842 18.81476063 4.93279847]

```

	Z	theta	nu	N
\count	1000000.000000	1000000.000000	1000000.000000	1000000.000000
mean	0.870055	10.001843	10.000098	4.999642
std	0.336243	5.774806	5.771397	2.581812
min	0.000000	0.000008	0.000020	1.000000
25%	1.000000	4.999020	4.999096	3.000000
50%	1.000000	9.997342	10.002818	5.000000

75%	1.000000	15.006256	14.988187	7.000000
max	1.000000	19.999982	19.999990	9.000000

```

count  1000000.000000
mean    4.994021
std     2.581949
min     1.000000
25%     3.000000
50%     5.000000
75%     7.000000
max     9.000000

```

Write Custom Data Loader

```
# %writefile data/dataloader.py
```

```

def split_t_x(df, target, source):
    # change from pandas dataframe format to a numpy
    # array of the specified types
    t = np.array(df[target])
    x = np.array(df[source])
    return t, x

def get_batch(x, batch_size):
    # the numpy function choice(length, number)
    # selects at random "batch_size" integers from
    # the range [0, length-1] corresponding to the
    # row indices.
    rows = np.random.choice(len(x), batch_size)
    batch_x = x[rows]
    # batch_x.T[-1] = np.random.uniform(0, 1, batch_size)
    return batch_x

def get_data_sets(simulate_data, batchsize):
    """write custom data generator because who wants to read pytorch's
    DataLoader source code
    (and its sometimes slow for some reason)"""
    # if simulate_data:
    #     Train_data_MLE_True = generate_training_data(Bprime=100000,
    MLE=True, save_data=False)

    # if SUBSAMPLE:

```



```

# data=load_df('data/TWO_PARAMETERS_TRAINING_DATA_1M.csv',
SUBSAMPLE=10000)#This is MLE DATA!
# else:
# data=load_df('data/TWO_PARAMETERS_TRAINING_DATA_1M.csv')
# data=load_df('data/TWO_PARAMETERS_TRAINING_DATA_1M.csv')
data =
pd.read_csv('data/two_parameters_theta_0_20_1000k_Examples_MLE_True.csv',
            # nrows=SAMPLE_SIZE,
            usecols=['Z','theta','nu','N','M']
            )
train_data, test_data = train_test_split(data,
                                         test_size=0.02)
#split the train data (0.8 of whole set) again into 0.8*0.8=0.64
of whole set
# train_data, valid_data = train_test_split(train_data,
test_size=0.2)
train_data = train_data.reset_index(drop=True)
# valid_data = valid_data.reset_index(drop=True)
test_data = test_data.reset_index(drop=True)

target='Z'
source = ['theta','nu','N','M']
train_t, train_x = split_t_x(train_data, target=target,
source=source)
# valid_t, valid_x = split_t_x(valid_data, target=target,
source=source)
test_t, test_x = split_t_x(test_data, target=target,
source=source)

def training_set_features():
    #start with an infinite loop, so that you can keep calling
    next (i.e. set = train_set(); set.next() ) until you run out of
    training examples
    while True:
        #get a random batch of the defined size
        batch_x = get_batch(train_x, batchsize)
        #print('batch_x', batch_x)
        #index of one of the items in our examples
        yield batch_x

def evaluation_set_features():
    #start with an infinite loop, so that you can keep calling
    next (i.e. set = train_set(); set.next() ) until you run out of
    training examples
    while True:
        batch_x = get_batch(test_x,batchsize)
        #index of one of the items in our examples
        yield batch_x

```

```

def training_set_targets():
    #start with an infinite loop, so that you can keep calling
    next (i.e. set = train_set(); set.next() ) until you run out of
    training examples
    while True:
        #get a random batch of the defined size
        batch_x = get_batch(train_t, batchsize)
        #print('batch_x', batch_x)
        #index of one of the items in our examples
        yield batch_x

```

```

def evaluation_set_targets():
    #start with an infinite loop, so that you can keep calling
    next (i.e. set = train_set(); set.next() ) until you run out of
    training examples
    while True:
        #get a random batch of the defined size
        batch_x = get_batch(test_t, batchsize)
        #print('batch_x', batch_x)
        #index of one of the items in our examples
        yield batch_x

```

```

return training_set_features, training_set_targets,
evaluation_set_features, evaluation_set_targets

```

```

training_set_features, training_set_targets, evaluation_set_features,
evaluation_set_targets = get_data_sets(simulate_data=True,
batchsize=300)

```

```

first_features_batch = next(training_set_features())
print('first features batch \n', first_features_batch[:5])
print('\nfirst features batch shape \n', first_features_batch.shape)

```

```

first features batch
[[19.52477661  6.51129404  2.          4.          ]
 [ 9.35131833 12.84672943  1.          6.          ]
 [ 0.61313754 13.7363275   7.          9.          ]
 [ 7.02718582 17.24384141  4.          7.          ]
 [14.6905525  16.97985822  5.          4.          ]]

```

```

first features batch shape
(300, 4)

```

```

def binsize(data_df_or_Bprime):
    """ decide on the number of bins in a histogram according to
    Struge's rule"""
    if isinstance(data_df_or_Bprime,pd.core.frame.DataFrame):
        bin_count = int(np.ceil(np.log2(data_df_or_Bprime.shape[0])) +
1)

```

```

elif isinstance(data_df_or_Bprime, int):
    bin_count = int(np.ceil(np.log2(data_df_or_Bprime)) + 1)
return bin_count

```

Plot the histogrammed function the histogrammed function

$h(\tilde{\theta}, v, N, M) = h(\theta_{min}, \theta_{max}, v, N, M)$ where $\tilde{\theta}$ means that it is simulated (inside the function).

```

def make_hist_data(Bprime,
                   thetamin, thetamax,
                   nu, N, M,
                   nbins,
                   MLE=True):

    theta = st.uniform.rvs(thetamin, thetamax, size=Bprime)
    n = st.poisson.rvs(theta + nu, size=Bprime)
    m = st.poisson.rvs(nu, size=Bprime)

    Z = (lambda_test(theta, n, m, MLE=MLE) <
         lambda_test(theta, N, M, MLE=MLE)).astype(np.int32)

    thetarange = (thetamin, thetamax)
    # bins = binsize(Bprime)

    # weighted histogram (count the number of ones per bin)
    y1, bb = np.histogram(theta,
                           bins=nbins,
                           range=thetarange,
                           weights=Z)

    # unweighted histogram (count number of ones and zeros per bin)
    yt, _ = np.histogram(theta,
                           bins=nbins,
                           range=thetarange)

    y = y1 / yt

    return y, bb

h, h_bins = make_hist_data(Bprime=1000,
                           thetamin=0, thetamax=20,
                           nu=3, N=1, M=1,
                           nbins=200, MLE=True)

```

h

```

array([0.25      , 0.28571429, 0.          , 0.33333333, 0.42857143,
        0.          , 0.33333333, 0.16666667, 0.16666667, 0.28571429,
        0.33333333, 0.25      , 0.71428571, 1.          , 0.66666667,
        0.75      , 0.71428571, 0.625      , 0.75      , 0.66666667,
        1.          , 0.875      , 0.8          , 1.          , 0.625      ,
        0.5          , 1.          , 0.81818182, 1.          , 0.8          ,


```

```
def plot_one_hist(Bprime, thetamin, thetamax, nu, N, M, MLE, nbins,
ax):
    counts, bins= make_hist_data(Bprime,
                                thetamin, thetamax,
                                nu, N, M,
                                nbins,
                                MLE)
    bin_centers = (bins[1:]+bins[:-1])/2
    ax.plot(bin_centers, counts, label= r'$\mathbf{h}$ Example', lw=3)
    ax.set_xlabel(r'$\theta$', fontsize=font_axes)
    ax.set_ylabel(r'$E(Z|\theta, \nu)$', fontsize=font_axes)
    ax.legend(loc='center right', fontsize=font_legend)

fig= plt.figure(); ax=fig.add_subplot()
```

#Example:

```
plot_one_hist(Bprime=100000, thetamin=0, thetamax=20,  
nu=3, N=1, M=1, MLE=True, nbins=200, ax=ax)
```



Plot the histogrammed approximations for the MLE vs non-MLE cases for a single value of ν

```
def plot_data_one_nu(Bprime, thetamin, thetamax, nu, D, MLE,  
    FONTSIZE=15,  
    func=None,  
    figsize=(10, 6)):
```

```
    # make room for 6 sub-plots
```

```
    fig, ax = plt.subplots(nrows=2,  
                           ncols=3,  
                           figsize=figsize)
```

```
    # padding
```

```
    plt.subplots_adjust(hspace=0.01)  
    plt.subplots_adjust(wspace=0.20)
```

```
    # use flatten() to convert a numpy array of  
    # shape (nrows, ncols) to a 1-d array.  
    ax = ax.flatten()
```

```
    for j, (N, M) in enumerate(D):
```

```
        y, bb = make_hist_data(Bprime,  
                                thetamin, thetamax,  
                                nu, N, M,
```

```

        nbins=200,
        MLE=True)

    ax[j].set_xlim(thetamin, thetamax-5)
    ax[j].set_ylim(0, 1)
    ax[j].set_xlabel(r'$\theta$', fontsize=FONTSIZE)
    ax[j].set_ylabel(r'$E(Z|\theta, \nu)$', fontsize=FONTSIZE)

    x = (bb[1:]+bb[:-1])/2
    ax[j].plot(x, y, 'b', lw=2, label='$\mathbf{h}$, MLE',
alpha=0.3)
    #h is histogram approximation

    y_nonMLE, bb_nonMLE = make_hist_data(Bprime,
        thetamin, thetamax,
        nu, N, M,
        nbins=200,
        MLE=False)

    x_nonMLE = (bb_nonMLE[1:]+bb_nonMLE[:-1])/2
    ax[j].plot(x_nonMLE, y_nonMLE, 'r', lw=2, label='$\mathbf{h}$,
non-MLE',alpha=0.3)

    if func:
        p, _ = func(nu, N, M)
        ax[j].plot(x, p, 'r', lw=2, label='f')
        #f is model approximation

        ax[j].grid(True, which="both", linestyle='-')
        ax[j].text(5.1, 0.42, r'$N, M = %d, %d$' % (N, M),
fontsize=font_legend-3
        # fontsize=FONTSIZE
        )

        ax[j].text(5.1, 0.30, r'$\nu = %5.1f$' % nu,
fontsize=font_legend-3
        # fontsize=FONTSIZE
        )

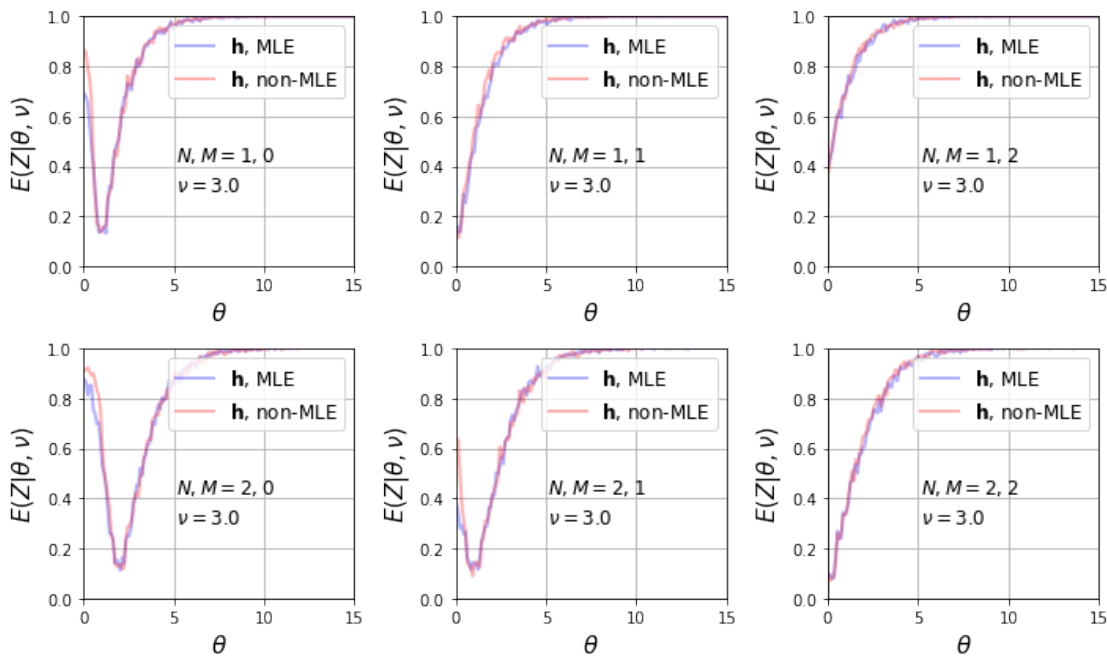
    ax[j].legend(loc='upper right',fontsize=font_legend-3)

# hide unused sub-plots
    for k in range(j+1, len(ax)):
        ax[k].set_visible(False)

    plt.tight_layout()
    plt.show()

```

```
D = [ (N, M) for N in range(1,3) for M in range(3)]
plot_data_one_nu(Bprime=100000, thetamin=0, thetamax=20, nu=3, D=D,
MLE=True)
```



Plot the histogrammed approximations for the MLE vs non-MLE cases for multiple values of ν , indicating the dependence on the nuisance parameter

```
def plot_data(Bprime, thetamin, thetamax, D, MLE,
              FONTSIZE=15,
              func=None,
              figsize=(15, 13)):
```

```
    # make room for 6 sub-plots
    fig, ax = plt.subplots(nrows=6,
                           ncols=2,
                           figsize=figsize)
```

```
    # padding
    plt.subplots_adjust(hspace=0.01)
    plt.subplots_adjust(wspace=0.20)
```

```
    # use flatten() to convert a numpy array of
    # shape (nrows, ncols) to a 1-d array.
    # ax = ax.flatten()
```

```
    for j, (N, M) in enumerate(D):
        NU1=1
        y, bb = make_hist_data(Bprime,
                                thetamin, thetamax,
                                NU1, N, M,
                                nbins=200,
```

```

MLE=True)

ax[j, 0].set_xlim(thetamin, thetamax-5)
ax[j, 0].set_ylim(0, 1.03)
ax[j, 0].set_xlabel(r'$\theta$', fontsize=FONTSIZE)
ax[j, 0].set_ylabel(r'$E(Z|\theta, \nu)$', fontsize=FONTSIZE)

x = (bb[1:]+bb[:-1])/2
ax[j, 0].plot(x, y, 'b', lw=2, label='$\mathbf{h}$, MLE',
alpha=0.45)
#h is histogram approximation

y_nonMLE, bb_nonMLE = make_hist_data(Bprime,
                                     thetamin, thetamax,
                                     NU1, N, M,
                                     nbins=200,
                                     MLE=False)

x_nonMLE = (bb_nonMLE[1:]+bb_nonMLE[:-1])/2
ax[j, 0].plot(x_nonMLE, y_nonMLE, 'r', lw=2, label='$\mathbf{h}$, non-MLE', alpha=0.45)

ax[j, 0].grid(True, which="both", linestyle='-')
ax[j, 0].text(6.1, 0.42, r'$N, M = %d, %d$' % (N, M),
             # fontsize=FONTSIZE
             )

ax[j, 0].text(6.1, 0.30, r'$\nu = %5.1f$' % NU1,
             # fontsize=FONTSIZE
             )

ax[j, 0].legend(loc='upper right', fontsize=font_legend-3)
#####define second nu value to see difference
NU2=9
y, bb = make_hist_data(Bprime,
                       thetamin, thetamax,
                       NU2, N, M,
                       nbins=200,
                       MLE=True)

ax[j, 1].set_xlim(thetamin, thetamax-5)
ax[j, 1].set_ylim(0, 1.03)
ax[j, 1].set_xlabel(r'$\theta$', fontsize=FONTSIZE)
ax[j, 1].set_ylabel(r'$E(Z|\theta, \nu)$', fontsize=FONTSIZE)

x = (bb[1:]+bb[:-1])/2
ax[j, 1].plot(x, y, 'b', lw=2, label='$\mathbf{h}$, MLE',
alpha=0.45)
#h is histogram approximation

```



```

y_nonMLE, bb_nonMLE = make_hist_data(Bprime,
                                     thetamin, thetamax,
                                     NU1, N, M,
                                     nbins=200,
                                     MLE=False)

x_nonMLE = (bb_nonMLE[1:]+bb_nonMLE[:-1])/2
ax[j, 1].plot(x_nonMLE, y_nonMLE, 'r', lw=2, label='$\mathbf{h}$, non-MLE', alpha=0.45)

ax[j, 1].grid(True, which="both", linestyle='-')
ax[j, 1].text(6.1, 0.42, r'$N, M = %d, %d$' % (N, M),
             # fontsize=FONTSIZE
             )

ax[j, 1].text(6.1, 0.30, r'$\nu = %5.1f$' % NU2,
             # fontsize=FONTSIZE
             )

ax[j, 1].legend(loc='upper right', fontsize=font_legend-3)

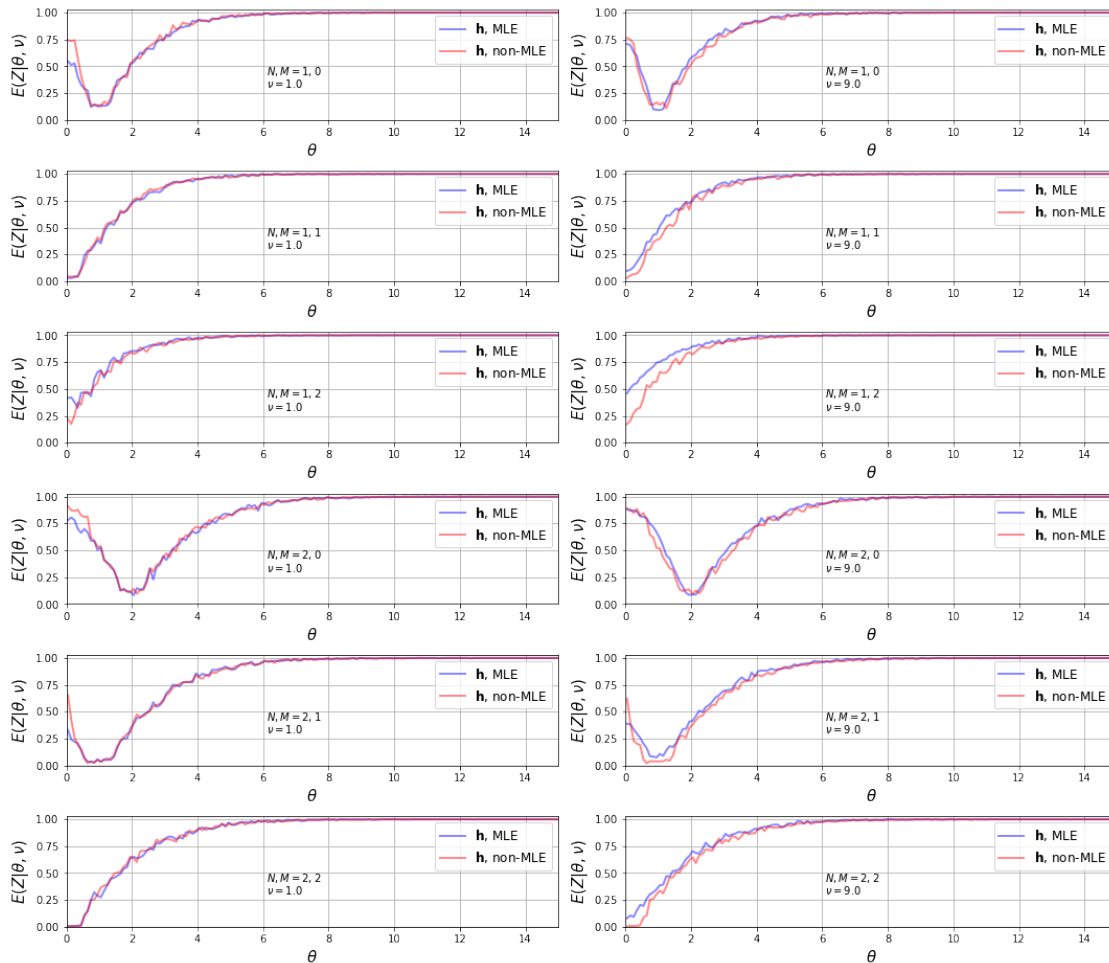
if func:
    p, _ = func(nu, N, M)
    ax[j, 0].plot(x, p, 'r', lw=2, label='f')
    # f is model approximation

# hide unused sub-plots
for k in range(j+1, len(ax)):
    ax[k].set_visible(False)

plt.tight_layout()
plt.show()

D = [ (N, M) for N in range(1,3) for M in range(3)]
plot_data(Bprime=100000, thetamin=0, thetamax=20, D=D, MLE=True)

```



ML

```
def getwholedata(MLE_or_nonMLE, valid=False):
    if MLE:
        data =
pd.read_csv('data/two_parameters_theta_0_20_1000k_Examples_MLE_True.csv',
            # nrows=SUBSAMPLE,
            usecols=['Z', 'theta', 'nu', 'N', 'M']
        )
    else:
        data =
pd.read_csv('data/two_parameters_theta_0_20_1000k_Examples_MLE_False.csv',
            # nrows=SUBSAMPLE,
            usecols=['Z', 'theta', 'nu', 'N', 'M']
        )
```

```

train_data, test_data = train_test_split(data, test_size=0.2)
#split the train data (0.8 of whole set) again into 0.8*0.8=0.64
of whole set

```

```

train_data = train_data.reset_index(drop=True)
test_data = test_data.reset_index(drop=True)

```

```

target='Z'
source = ['theta', 'nu', 'N', 'M']

```

```

train_t, train_x = split_t_x(train_data, target=target,
source=source)
test_t, test_x = split_t_x(test_data, target=target,
source=source)
print('train_t shape = ', train_t.shape, '\n')
print('train_x shape = ', train_x.shape, '\n')

```

```

if valid:
    #if you want to also make a validation data set
    train_data, valid_data = train_test_split(train_data,
test_size=0.2)
    valid_data = valid_data.reset_index(drop=True)
    valid_t, valid_x = split_t_x(valid_data, target=target,
source=source)

```

```

return train_t, train_x, test_t, test_x

```

```

# n_iterations=50

```

Define Model, which will approximate the expectation value above

```

class Model(nn.Module):

```

```

    def __init__(self, n_inputs=4, n_nodes=20, n_layers=5):

```

```

        # call constructor of base (or super, or parent) class
        super(Model, self).__init__()

```

```

        self.layers = []

```

```

        # create input layer
        self.layer0 = nn.Linear(n_inputs, n_nodes)
        self.layers.append(self.layer0)

```

```

        # create "hidden" layers
        for l in range(1, n_layers):
            cmd = 'self.layer%d = nn.Linear(%d, %d)' % \

```

```

        (l, n_nodes, n_nodes)
        exec(cmd)
        cmd = 'self.layers.append(self.layer%d)' % l
        exec(cmd)

    # create output layer
    cmd = 'self.layer%d = nn.Linear(%d, 1)' % (n_layers, n_nodes)
    exec(cmd)
    cmd = 'self.layers.append(self.layer%d)' % n_layers
    exec(cmd)

    # define (required) method to compute output of network
    def forward(self, x):
        y = x
        for layer in self.layers[:-1]:
            y = layer(y)
            y = torch.relu(y)
        y = self.layers[-1](y)
        y = torch.sigmoid(y)
        return y

model = Model()
print(model)

Model(
  (layer0): Linear(in_features=4, out_features=20, bias=True)
  (layer1): Linear(in_features=20, out_features=20, bias=True)
  (layer2): Linear(in_features=20, out_features=20, bias=True)
  (layer3): Linear(in_features=20, out_features=20, bias=True)
  (layer4): Linear(in_features=20, out_features=20, bias=True)
  (layer5): Linear(in_features=20, out_features=1, bias=True)
)

def average_quadratic_loss(f, t, x):
    # f and t must be of the same shape
    return torch.mean((f - t)**2)
def average_loss(f, t):
    # f and t must be of the same shape
    return torch.mean((f - t)**2)

def validate(model, avloss, inputs, targets):
    # make sure we set evaluation mode so that any training specific
    # operations are disabled.
    model.eval() # evaluation mode

    with torch.no_grad(): # no need to compute gradients wrt. x and t
        x = torch.from_numpy(inputs).float()
        t = torch.from_numpy(targets).float()
        # remember to reshape!
        o = model(x).reshape(t.shape)

```

```

    return avloss(o, t, x)

def get_features_training_batch(x, t, batch_size):
    # the numpy function choice(length, number)
    # selects at random "batch_size" integers from
    # the range [0, length-1] corresponding to the
    # row indices.
    rows = np.random.choice(len(x), batch_size)
    batch_x = x[rows]
    batch_t = t[rows]
    # batch_x.T[-1] = np.random.uniform(0, 1, batch_size)
    return (batch_x, batch_t)

def train(model, optimizer, avloss,
          batch_size,
          n_iterations, traces,
          step, window, MLE):

    # to keep track of average losses
    xx, yy_t, yy_v, yy_v_avg = traces

    n = len(valid_x)

    if MLE:
        train_t, train_x, test_t, test_x =
getwholedata(MLE_or_nonMLE=True, valid=False)
    else:
        train_t, train_x, test_t, test_x =
getwholedata(MLE_or_nonMLE=False, valid=False)

    print('Iteration vs average loss')
    print("%10s\t%10s\t%10s" % \
          ('iteration', 'train-set', 'valid-set'))

    # training_set_features, training_set_targets,
    evaluation_set_features, evaluation_set_targets =
get_data_sets(simulate_data=False, batchsize=batch_size)

    for ii in range(n_iterations):

        # set mode to training so that training specific
        # operations such as dropout are enabled.

        model.train()

        # get a random sample (a batch) of data (as numpy arrays)

```

```

    #Harrison-like Loader
    batch_x, batch_t = get_features_training_batch(train_x,
train_t, batch_size)

    #Or Ali's Loader
    # batch_x, batch_t = next(training_set_features()),
next(training_set_targets())
    # batch_x_eval, batch_t_eval =
next(evaluation_set_features()), next(evaluation_set_targets())

    with torch.no_grad(): # no need to compute gradients
        # wrt. x and t
        x = torch.from_numpy(batch_x).float()
        t = torch.from_numpy(batch_t).float()

    outputs = model(x).reshape(t.shape)

    # compute a noisy approximation to the average loss
    empirical_risk = avloss(outputs, t, x)

    # use automatic differentiation to compute a
# noisy approximation of the local gradient
    optimizer.zero_grad()      # clear previous gradients
    empirical_risk.backward()    # compute gradients

    # finally, advance one step in the direction of steepest
# descent, using the noisy local gradient.
    optimizer.step()           # move one step

    if ii % step == 0:

        #using Harrison-like loader
        acc_t = validate(model, avloss, train_x[:n], train_t[:n])
        acc_v = validate(model, avloss, test_x[:n], test_t[:n])

        #using Ali's loader
        # acc_t = validate(model, avloss, batch_x, batch_t)
        # acc_v = validate(model, avloss, batch_x_eval,
batch_t_eval)

    yy_t.append(acc_t)
    yy_v.append(acc_v)

    # compute running average for validation data
    len_yy_v = len(yy_v)
    if len_yy_v < window:
        yy_v_avg.append( yy_v[-1] )

```

```

        elif len_yy_v == window:
            yy_v_avg.append( sum(yy_v) / window )
        else:
            acc_v_avg = yy_v_avg[-1] * window
            acc_v_avg += yy_v[-1] - yy_v[-window-1]
            yy_v_avg.append(acc_v_avg / window)

    if len(xx) < 1:
        xx.append(0)
        print("%10d\t%10.6f\t%10.6f" % \
              (xx[-1], yy_t[-1], yy_v[-1]))
    else:
        xx.append(xx[-1] + step)

        print("\r%10d\t%10.6f\t%10.6f\t%10.6f" % \
              (xx[-1], yy_t[-1], yy_v[-1], yy_v_avg[-1]),
              end='')

print()
return (xx, yy_t, yy_v, yy_v_avg)

def plot_average_loss(traces, ftsize=18, save_loss_plots=False):

    xx, yy_t, yy_v, yy_v_avg = traces

    # create an empty figure
    fig = plt.figure(figsize=(6, 4.5))
    fig.tight_layout()

    # add a subplot to it
    nrows, ncols, index = 1,1,1
    ax = fig.add_subplot(nrows,ncols,index)

    ax.set_title("Average loss")

    ax.plot(xx, yy_t, 'b', lw=2, label='Training')
    ax.plot(xx, yy_v, 'r', lw=2, label='Validation')
    #ax.plot(xx, yy_v_avg, 'g', lw=2, label='Running average')

    ax.set_xlabel('Iterations', fontsize=ftsiz)
    ax.set_ylabel('average loss', fontsize=ftsiz)
    ax.set_xscale('log')
    ax.set_yscale('log')
    ax.grid(True, which="both", linestyle='-')
    ax.legend(loc='upper right')
    if save_loss_plots:
        plt.savefig('images/loss_curves/IQN_'+N+T+'_Consecutive_2.png')
        print('\nloss curve saved in
images/loss_curves/IQN_'+N+target+'_Consecutive.png')

```

```
# if show_loss_plots:
plt.show()
```

Define my regularized regression model. Since the values are on the same scales, it is not necessary to include batchnormalization or to normalize the data

```
class RegularizedRegressionModel(nn.Module):
    #inherit from the super class
    def __init__(self, nfeatures, ntargets, nlayers, hidden_size,
dropout):
        super().__init__()
        layers = []
        for i in range(nlayers):
            if len(layers) == 0:
                #initial layer has to have size of input features as
its input layer
                #its output layer can have any size but it must match
the size of the input layer of the next linear layer
                #here we choose its output layer as the hidden size
(fully connected)
                layers.append(nn.Linear(nfeatures, hidden_size))
                #batch normalization
                # layers.append(nn.BatchNorm1d(hidden_size))
                #Dropout seems to worsen model performance
                layers.append(nn.Dropout(dropout))
                #ReLU activation
                layers.append(nn.ReLU())
            else:
                #if this is not the first layer (we dont have layers)
                layers.append(nn.Linear(hidden_size, hidden_size))
                # layers.append(nn.BatchNorm1d(hidden_size))
                #Dropout seems to worsen model performance
                layers.append(nn.Dropout(dropout))
                layers.append(nn.ReLU())
                #output layer:
        layers.append(nn.Linear(hidden_size, ntargets))

        # ONLY IF ITS A CLASSIFICATION, ADD SIGMOID
        layers.append(nn.Sigmoid())
        #we have defined sequential model using the layers in
oulist
        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)
```


Make a hyperparameter Tuning Workflow

Use Optuna ([axriv:1907.10902](https://arxiv.org/abs/1907.10902)) for hyperparameter tuning. The search space for the hyperparameters that I'm tuning is defined in the params dictionary:

```
params = {
    "nlayers": trial.suggest_int("nlayers",1,13),
    "hidden_size": trial.suggest_int("hidden_size", 2, 130),
    "dropout": trial.suggest_float("dropout", 0.1,0.5),
    "optimizer_name" : trial.suggest_categorical("optimizer_name",
["Adam", "RMSprop"]),
    "learning_rate": trial.suggest_float("learning_rate", 1e-5, 1e-2),
    "batch_size": trial.suggest_int("batch_size", 1000, 10000)
}
```

class Engine:

```
    """loss, training and evaluation"""
    def __init__(self, model, optimizer, batch_size):
        #, device):
        self.model = model
        #self.device= device
        self.optimizer = optimizer
        self.batch_size=batch_size

    #the loss function returns the loss function. It is a static
method so it doesn't need self
    # @staticmethod
    # def loss_fun(targets, outputs):
    #     tau = torch.rand(outputs.shape)
    #     return torch.mean(torch.where(targets >= outputs,
    #                                     tau * (targets - outputs),
    #                                     (1 - tau)*(outputs -
targets)))

    #     This loss combines a Sigmoid layer and the BCELoss in one single
class. This version is more numerically stable than using a plain
Sigmoid followed by a BCELoss as,
    #     by combining the operations into one layer

    def train(self, x, t):
        """the training function: takes the training dataloader"""
        self.model.train()
        final_loss = 0
        for iteration in range(n_iterations):
            self.optimizer.zero_grad()
            batch_x, batch_t = get_features_training_batch(x, t,
self.batch_size)#x and t are train_x and train_t
```

```

        # with torch.no_grad():
        inputs=torch.from_numpy(batch_x).float()
        targets=torch.from_numpy(batch_t).float()
        outputs = self.model(inputs)
        loss = average_quadratic_loss(outputs, targets, inputs)
        loss.backward()
        self.optimizer.step()
        final_loss += loss.item()

    return final_loss / self.batch_size

def evaluate(self, x, t):
    """the training function: takes the training dataloader"""
    self.model.eval()
    final_loss = 0
    for iteration in range(n_iterations):
        batch_x, batch_t = get_features_training_batch(x, t,
self.batch_size)#x and t are train_x and train_t

        # with torch.no_grad():
        inputs=torch.from_numpy(batch_x).float()
        targets=torch.from_numpy(batch_t).float()
        outputs = self.model(inputs)
        loss =average_quadratic_loss(outputs, targets, inputs)
        final_loss += loss.item()
    return final_loss / self.batch_size

```

EPOCHS=1

```

def run_train(params, save_model=False):
    """For tuning the parameters"""

    model = RegularizedRegressionModel(
        nfeatures=train_x.shape[1],
        ntargets=1,
        nlayers=params["nlayers"],
        hidden_size=params["hidden_size"],
        dropout=params["dropout"]
    )
    # print(model)

    learning_rate= params["learning_rate"]
    optimizer_name = params["optimizer_name"]

    # optimizer = torch.optim.Adam(model.parameters(),
lr=params["learning_rate"])

```

```

optimizer = getattr(torch.optim, optimizer_name)
(model.parameters(), lr=learning_rate)

eng=Engine(model, optimizer, batch_size=params["batch_size"])
best_loss = np.inf
early_stopping_iter=10
early_stopping_coutner=0

for epoch in range(EPOCHS):
    train_loss = eng.train(train_x, train_t)
    valid_loss=eng.evaluate(test_x, test_t)

    print(f"{epoch} \t {train_loss} \t {valid_loss}")
    if valid_loss<best_loss:
        best_loss=valid_loss
        if save_model:
            model.save(model.state_dict(), "model_m.bin")
    else:
        early_stopping_coutner+=1
    if early_stopping_coutner > early_stopping_iter:
        break
return best_loss

# run_train()

def objective(trial):
    params = {
        "nlayers": trial.suggest_int("nlayers",1,13),
        "hidden_size": trial.suggest_int("hidden_size", 2, 130),
        "dropout": trial.suggest_float("dropout", 0.1,0.5),
        "optimizer_name" : trial.suggest_categorical("optimizer_name",
["Adam", "RMSprop"]),
        "learning_rate": trial.suggest_float("learning_rate", 1e-5, 1e-
2),
        "batch_size": trial.suggest_int("batch_size", 1000, 10000)

    }
    # all_losses=[]

    temp_loss = run_train(params,save_model=False)
    # all_losses.append(temp_loss)
    return temp_loss

def tune_hyperparameters():
    print('Getting best hyperparameters')
    study=optuna.create_study(direction="minimize")
    study.optimize(objective, n_trials=10)
    best_trial = study.best_trial
    print('best model parameters', best_trial.params)

```

```

best_params=best_trial.params#this is a dictionary
filename='best_params/best_params_Test_Trials.csv'
param_df=pd.DataFrame({
    'n_layers':best_params["nlayers"],
    'hidden_size':best_params["hidden_size"],
    'dropout':best_params["dropout"],

    'optimizer_name':best_params["optimizer_name"],
    'learning_rate':
best_params["learning_rate"],
    'batch_size':best_params["batch_size"] },
    index=[0]
)

param_df.to_csv(filename)

```

Don't run the one cell below, unless you want to tune!

```
tune_hyperparameters()
```

Load the dictionary of the best hyperparameters that was saved from our hyperparameter tuning workflow, and retrieve the values

```

BEST_PARAMS = pd.read_csv('best_params/best_params_Test_Trials.csv')
print(BEST_PARAMS)

```

```

n_layers = int(BEST_PARAMS["n_layers"])
hidden_size = int(BEST_PARAMS["hidden_size"])
dropout = float(BEST_PARAMS["dropout"])
optimizer_name = BEST_PARAMS["optimizer_name"].to_string().split()[1]
learning_rate = float(BEST_PARAMS["learning_rate"])
batch_size = int(BEST_PARAMS["batch_size"])

```

```

    Unnamed: 0  n_layers  hidden_size  dropout  optimizer_name
learning_rate  \
0              0          4          11  0.13208          RMSprop
0.006398

```

```

    batch_size
0          1000

```

Define network node shapes, parameters, and training data

```

BATCHSIZE=batch_size
training_set_features, training_set_targets, evaluation_set_features,
evaluation_set_targets = \
get_data_sets(simulate_data=False, batchsize=BATCHSIZE)

```

```

sample_x=next(training_set_features())#this is just to get the
dimenstions of one batch
sample_y=next(training_set_targets())
#(batchsize,5) for mass

```

```
print('sample x shape', sample_x.shape)
print('sample t shape', sample_y.shape)
```

```
n_features = sample_x.shape[1]
print('\n')
```

```
model = RegularizedRegressionModel(
    nfeatures=sample_x.shape[1],
    ntargets=1,
    nlayers=n_layers,
    hidden_size=hidden_size,
    dropout=dropout
)
```

```
print(model)
```

```
sample x shape (1000, 4)
sample t shape (1000,)
```

```
RegularizedRegressionModel(
(model): Sequential(
  (0): Linear(in_features=4, out_features=11, bias=True)
  (1): Dropout(p=0.1320798105984151, inplace=False)
  (2): ReLU()
  (3): Linear(in_features=11, out_features=11, bias=True)
  (4): Dropout(p=0.1320798105984151, inplace=False)
  (5): ReLU()
  (6): Linear(in_features=11, out_features=11, bias=True)
  (7): Dropout(p=0.1320798105984151, inplace=False)
  (8): ReLU()
  (9): Linear(in_features=11, out_features=11, bias=True)
  (10): Dropout(p=0.1320798105984151, inplace=False)
  (11): ReLU()
  (12): Linear(in_features=11, out_features=1, bias=True)
  (13): Sigmoid()
)
```

Chose whose model (Ali or Harrison) and paramers and train

```
def initiate_whose_model(Ali_or_Harrison, MLE):
    whose_model='Ali'

    if whose_model=='Harrison':
        n_layers=5
        hidden_size=5
        dropout=0
        learning_rate=int(1e-3)
```

```

        batch_size=64
        optimizer = torch.optim.Adam(model.parameters()),
lr=int(1e-3))
        model=Model()
        elif whose_model=='Ali':
            n_layers = int(BEST_PARAMS["n_layers"])
            hidden_size = int(BEST_PARAMS["hidden_size"])
            dropout = float(BEST_PARAMS["dropout"])
            optimizer_name =
BEST_PARAMS["optimizer_name"].to_string().split()[1]
            learning_rate = float(BEST_PARAMS["learning_rate"])
            batch_size = int(BEST_PARAMS["batch_size"])
            model = RegularizedRegressionModel(
                nfeatures=sample_x.shape[1],
                ntargets=1,
                nlayers=n_layers,
                hidden_size=hidden_size,
                dropout=dropout
            )
            optimizer = getattr(torch.optim, str(optimizer_name) )
            (model.parameters(), lr=learning_rate)

        return n_layers, hidden_size, dropout, optimizer_name,
learning_rate, batch_size, model, optimizer

n_layers, hidden_size, dropout, optimizer_name, learning_rate,
batch_size, model_MLE, optimizer_MLE = initiate_whose_model('Ali',
MLE=True)
print(optimizer_MLE)
print('\n\n')
print(model_MLE)

#also initiate non-MLE model
n_layers, hidden_size, dropout, optimizer_name, learning_rate,
batch_size, model_nonMLE, optimizer_nonMLE =
initiate_whose_model('Ali', MLE=False)

RMSprop (
Parameter Group 0
    alpha: 0.99
    centered: False
    eps: 1e-08
    lr: 0.0063975512794992
    momentum: 0
    weight_decay: 0
)

```

```

RegularizedRegressionModel(

```

```
(model): Sequential(
  (0): Linear(in_features=4, out_features=11, bias=True)
  (1): Dropout(p=0.1320798105984151, inplace=False)
  (2): ReLU()
  (3): Linear(in_features=11, out_features=11, bias=True)
  (4): Dropout(p=0.1320798105984151, inplace=False)
  (5): ReLU()
  (6): Linear(in_features=11, out_features=11, bias=True)
  (7): Dropout(p=0.1320798105984151, inplace=False)
  (8): ReLU()
  (9): Linear(in_features=11, out_features=11, bias=True)
  (10): Dropout(p=0.1320798105984151, inplace=False)
  (11): ReLU()
  (12): Linear(in_features=11, out_features=1, bias=True)
  (13): Sigmoid()
)
```

Train MLE model

BATCHSIZE=batch_size

traces_MLE = ([], [], [], [])

traces_step = 200

n_iterations=10000

```
traces_MLE= train(model=model_MLE,
                  optimizer=optimizer_MLE,
                  avloss=average_quadratic_loss,
                  batch_size=BATCHSIZE,
                  n_iterations=n_iterations,
                  traces=traces_MLE,
                  step=traces_step,
                  window=200,
                  MLE=True)
```

train_t shape = (800000,)

train_x shape = (800000, 4)

Iteration vs average loss

iteration	train-set	valid-set	
0	0.134987	0.134437	
9800	0.058935	0.058109	0.058109

n_iterations=10000

BATCHSIZE=500


traces= train(model=model, optimizer=optimizer,

avloss=average_quadratic_loss,

batch_size=BATCHSIZE,

```
#         n_iterations=n_iterations, traces=traces,
#         step=traces_step, window=100)

plot_average_loss(traces_MLE)
```



Train non-MLE model

```
BATCHSIZE=batch_size
traces_nonMLE = ([], [], [], [])
traces_step = 200

n_iterations=200000
traces_nonMLE= train(model=model_nonMLE,
                    optimizer=optimizer_nonMLE,
                    avloss=average_quadratic_loss,
                    batch_size=BATCHSIZE,
                    n_iterations=n_iterations,
                    traces=traces_nonMLE,
                    step=traces_step,
                    window=200,
                    MLE=False)
plot_average_loss(traces_nonMLE)


train_t shape = (800000,)
```



```
train_x shape = (800000, 4)
```

Iteration vs average loss

iteration	train-set	valid-set	
0	0.116824	0.115104	
199800	0.058369	0.058304	0.059388



Make sure the train df has the same ranges as the data you want to generate for evaluation

```
def load_train_df(MLE):  
    """ returns the dataframe, can be used if the dataframe is saved  
    in csv format  
    of if it is already in dataframe format (e.g. generated in this  
    notebook). """  
    # SUBSAMPLE=int(1e5)  
    # if isinstance(df_name,str):  
    if MLE:  
        train_df =  
pd.read_csv('data/two_parameters_theta_0_20_1000k_Examples_MLE_True.csv',  
             # nrow=SUBSAMPLE,  
             usecols=['Z','theta','nu', 'N', 'M']  
             )  
    else:  
        train_df =  
pd.read_csv('data/two_parameters_theta_0_20_1000k_Examples_MLE_False.c
```

```
sv',
        # n_rows=SUBSAMPLE,
        usecols=['Z', 'theta', 'nu', 'N', 'M']
    )
    return train_df
```

```
train_df_MLE = load_train_df(MLE=True)
train_df_MLE.describe()
```

	Z	theta	nu	N
\count	1000000.000000	1000000.000000	1000000.000000	1000000.000000
mean	0.870055	10.001843	10.000098	4.999642
std	0.336243	5.774806	5.771397	2.581812
min	0.000000	0.000008	0.000020	1.000000
25%	1.000000	4.999020	4.999096	3.000000
50%	1.000000	9.997342	10.002818	5.000000
75%	1.000000	15.006256	14.988187	7.000000
max	1.000000	19.999982	19.999990	9.000000

	M
count	1000000.000000
mean	4.994021
std	2.581949
min	1.000000
25%	3.000000
50%	5.000000
75%	7.000000
max	9.000000

Make "on-the-fly" generated evaluation data

Note that this is one advantage of LFI, where one can always generate more synthetic data (for training as well as evaluation), whereas in traditional ML, the training and evaluation data sets are fixed. Here, we generate binned θ with the same ranges as those of the training set, and constants for $\{v, N, M\}$.

```
def make_eval_data(Bprime, train_df, nu, N, M, nbins):
    #if MLE true, load the model that was trained on MLE data and vice versa
    # N, M = D
```

```

# nbins=NBINS
# thetamin,thetamax=0,20
thetamin=train_df['theta'].min()
thetamax=train_df['theta'].max()
thetastep = (thetamax-thetamin) / nbins
bb      = np.arange(thetamin, thetamax+thetastep, thetastep)#this is
just making a vector of thetas
X      = (bb[1:] + bb[:-1])/2
tensor = torch.Tensor([[x, nu, N, M] for x in X])
return tensor, X

```

Look at an example of a "on-the-fly" generated evaluation data

```

eval_data_example, eval_bins_example
=make_eval_data(Bprime=100,train_df=train_df_MLE,nu=3, N=1, M=3,
nbins=300)
print(eval_data_example[:5])

```

```

tensor([[0.0333, 3.0000, 1.0000, 3.0000],
        [0.1000, 3.0000, 1.0000, 3.0000],
        [0.1667, 3.0000, 1.0000, 3.0000],
        [0.2333, 3.0000, 1.0000, 3.0000],
        [0.3000, 3.0000, 1.0000, 3.0000]])

```

```
eval_data_example.shape
```

```
torch.Size([300, 4])
```

Evaluate Trained model at this generated data

```

def usemodel(Bprime, train_df, nu, N,M, MLE, nbins):
    #make evaluation data at those fixed nu, N, M values
    eval_data, eval_bins =make_eval_data(Bprime,train_df,nu, N,M,
nbins)#eval data is independent of MLE, since its just constants witha
theta variable

```

```

    # if MLE==True:
    #     model=model
    #else load the model trained on non-MLE data
    # PATH='models/MLE_TRUE_Regressor_200.0K_training_iter.pt'
    if MLE:
        PATH= 'models/MLE_TRUE_Regressor_200.0K_training_iter.pt'
    else:
        PATH= 'models/MLE_False_Regressor_200.0K_training_iter.pt'
    n_layers = int(BEST_PARAMS["n_layers"])
    hidden_size = int(BEST_PARAMS["hidden_size"])
    dropout = float(BEST_PARAMS["dropout"])
    optimizer_name = BEST_PARAMS["optimizer_name"].to_string().split()
[1] learning_rate = float(BEST_PARAMS["learning_rate"])
    batch_size = int(BEST_PARAMS["batch_size"])
    model = RegularizedRegressionModel(
        nfeatures=train_x.shape[1],

```

```

        ntargets=1,
        nlayers=n_layers,
        hidden_size=hidden_size,
        dropout=dropout
    )
    model.load_state_dict(torch.load(PATH) )
    model.eval()
    return model(eval_data).detach().numpy(), eval_bins

phat_MLE, phatbins_MLE = usemodel(Bprime=1000,train_df=train_df,nu=3,
N=2, M=3, MLE=True, nbins=200)
print(phat[:5])

```

```

[[0.3516053 ]
 [0.37797767]
 [0.40793678]
 [0.44479653]
 [0.4591592 ]]


```

```

plt.plot(phatbins_MLE, phat_MLE, label=r'$\mathbf{f}$ MLE');
plt.legend()

```

<matplotlib.legend.Legend at 0x7f1e92f43e10>



```

phat_nonMLE, phatbins_nonMLE =
usemodel(Bprime=1000,train_df=train_df,nu=3, N=2, M=3, MLE=False,
nbins=200)
print(phat[:5])
plt.plot(phatbins_nonMLE, phat_nonMLE, label=r'$\mathbf{f}$ non-MLE');
plt.legend()


```

```

[[0.3516053 ]
 [0.37797767]
 [0.40793678]
 [0.44479653]
 [0.4591592  ]]

```

<matplotlib.legend.Legend at 0x7f1e92e22990>



SAVE TRAINED MODEL (if it's good)

```

def save_model(MLE):
    if MLE:
        model = model_MLE
    else:
        model = model_nonMLE
    PATH='models/MLE_%s_Regressor_%sK_training_iter.pt' % ( str(MLE),
str(n_iterations/1000) )
    torch.save(model.state_dict(), PATH)

```

```
save_model(MLE=False)
```

This is how you load a trained model

```

#load
PATH='models/MLE_TRUE_Regressor_20.0K_training_iter.pt'
n_layers = int(BEST_PARAMS["n_layers"])
hidden_size = int(BEST_PARAMS["hidden_size"])
dropout = float(BEST_PARAMS["dropout"])
optimizer_name = BEST_PARAMS["optimizer_name"].to_string().split()[1]
learning_rate = float(BEST_PARAMS["learning_rate"])
batch_size = int(BEST_PARAMS["batch_size"])
model = RegularizedRegressionModel(

```

```

        nfeatures=train_x.shape[1],
        ntargets=1,
        nlayers=n_layers,
        hidden_size=hidden_size,
        dropout=dropout
    )
model.load_state_dict(torch.load(PATH) )
#OR
#model=torch.load(PATH)#BUT HERE IT WILL BE A DICT (CANT BE EVALUATED
RIGHT AWAY) DISCOURAGED!
model.eval()
print(model)

```

```

RegularizedRegressionModel(
  (model): Sequential(
    (0): Linear(in_features=4, out_features=11, bias=True)
    (1): Dropout(p=0.1320798105984151, inplace=False)
    (2): ReLU()
    (3): Linear(in_features=11, out_features=11, bias=True)
    (4): Dropout(p=0.1320798105984151, inplace=False)
    (5): ReLU()
    (6): Linear(in_features=11, out_features=11, bias=True)
    (7): Dropout(p=0.1320798105984151, inplace=False)
    (8): ReLU()
    (9): Linear(in_features=11, out_features=11, bias=True)
    (10): Dropout(p=0.1320798105984151, inplace=False)
    (11): ReLU()
    (12): Linear(in_features=11, out_features=1, bias=True)
    (13): Sigmoid()
  )
)

```

You could also evaluate the trained model on the validation data

```

# nbins=100
# thetamin=train_df['theta'].min()
# thetamax=train_df['theta'].max()
# thetastep = (thetamax-thetamin) / nbins
# bb      = np.arange(thetamin, thetamax+thetastep, thetastep)#this is
just making a vector of thetas
# X      = (bb[1:] + bb[:-1])/2
eval_data=torch.Tensor(valid_x)
model.eval()
phat=model(eval_data).detach().numpy()

```

phat

```

array([[0.8695913],
       [0.8695913],
       [0.8695913],
       ...,
       [0.8695913],

```

```

        [0.8695913],
        [0.8695913]], dtype=float32)

def plot_data_one_nu_with_model(Bprime, thetamin, thetamax, nu, D,
MLE,
                                NBINS,
                                FONTSIZE=15,
                                func=None,
                                figsize=(10, 6), save_image=False):

    # make room for 6 sub-plots
    fig, ax = plt.subplots(nrows=2,
                            ncols=3,
                            figsize=figsize)

    # padding
    plt.subplots_adjust(hspace=0.01)
    plt.subplots_adjust(wspace=0.20)

    # use flatten() to convert a numpy array of
    # shape (nrows, ncols) to a 1-d array.
    ax = ax.flatten()

    for j, (N, M) in enumerate(D):

        y, bb = make_hist_data(Bprime,
                                thetamin, thetamax,
                                nu, N, M,
                                nbins=NBINS,
                                MLE=True)

        ax[j].set_xlim(thetamin-0.5, thetamax-5)
        ax[j].set_ylim(0, 1.03)
        ax[j].set_xlabel(r'\mathbf{\theta}', fontsize=FONTSIZE-3)
        ax[j].set_ylabel(r'\mathbf{E(Z|\theta, \nu)}',
        fontsize=FONTSIZE-3)

        x = (bb[1:]+bb[:-1])/2
        ax[j].plot(x, y, 'b', lw=2, label='\mathbf{h}$ MLE',
alpha=0.4)
        #h is histogram approximation

        y_nonMLE, bb_nonMLE = make_hist_data(Bprime,
                                                thetamin, thetamax,
                                                nu, N, M,
                                                nbins=NBINS,
                                                MLE=False)

        x_nonMLE = (bb_nonMLE[1:]+bb_nonMLE[:-1])/2

```

```

        ax[j].plot(x_nonMLE, y_nonMLE, 'r', lw=2, label='$\mathbf{h}$
non-MLE', alpha=0.4)

        if func:
            train_df_MLE = load_train_df(MLE=True)
            train_df_nonMLE = load_train_df(MLE=False)

            f_MLE, f_bins_MLE = func(Bprime, train_df_MLE, nu, N, M,
MLE=True, nbins=NBINS)
            ax[j].plot(x, f_MLE, color='g', lw=2, label='$\mathbf{f}$
MLE', alpha=0.4)
            # f is model approximation

            f_nonMLE, f_bins_nonMLE = func(Bprime, train_df_nonMLE,
nu, N, M, MLE=False, nbins=NBINS)
            ax[j].plot(x, f_nonMLE, color='c', lw=2, label='$\
mathbf{f}$ non-MLE', alpha=0.4)

            ax[j].grid(True, which="both", linestyle='-')
            ax[j].text(3.1, 0.42, r'$N, M = %d, %d$' % (N, M),
fontsize=font_legend-3
                        # fontsize=FONTSIZE
                        )

            ax[j].text(3.1, 0.30, r'$\nu = %5.1f$' % nu,
fontsize=font_legend-3
                        # fontsize=FONTSIZE
                        )

            ax[j].legend(loc='upper right', fontsize=font_legend-3)

        # hide unused sub-plots
        for k in range(j+1, len(ax)):
            ax[k].set_visible(False)

    plt.tight_layout()
    if save_image:
        plt.savefig('images/h_MLE_nonMLE_f_MLE_f_nonMLE_one_nu.png')
    plt.show()

```

Compare the histogrammed function $h(\theta, \nu, N, M)$ to the ML prediction function $f(\theta, \nu, N, M)$ (which is trained to regress Z).

```

def make_D(train_df):
    Nmin = train_df['N'].min()
    Nmax = train_df['N'].max()
    Mmin = train_df['M'].min()
    Mmax = train_df['M'].max()
    D = [ (N, M) for N in range(Nmin, Nmax) for M in range(Mmin,

```



```

Mmax)]
    return D[:6]

D = make_D(train_df)
D

[(1, 1), (1, 2), (1, 3), (1, 4), (1, 5), (1, 6)]

thetamin = train_df['theta'].min()
thetamax = train_df['theta'].max()
plot_data_one_nu_with_model(Bprime=100000, thetamin=thetamin,
thetamax=thetamax,
    nu=3, D=D, MLE=True, NBINS=200, func=usemodel,
save_image=True)

```

